

LAPENNA Program

Q&A 2

Kwai Wong, Stan Tomov
Julian Halloy, Stephen Qiu, Eric Zhao

University of Tennessee, Knoxville

Sept. 29, 2021

Acknowledgements:

- Support from NSF, UTK, JICS, ICL, NICS
- LAPENNA, www.jics.utk.edu/lapenna, NSF award #202409
- www.icl.utk.edu, cfdlab.utk.edu, www.xsede.org,
www.jics.utk.edu/recsem-reu,
- MagmaDNN is a project grown out from the RECSEM REU Summer program supported under NSF award #1659502
- Source code: www.bitbucket.org/icl/magmadnn
- www.bitbucket.org/cfdl/opendnnwheel



INNOVATIVE
COMPUTING LABORATORY

THE UNIVERSITY OF
TENNESSEE
KNOXVILLE

JICS
Joint Institute for
Computational Sciences
ORNL
Computational Sciences

OAK RIDGE
National Laboratory

Exercises

Multilayer Perceptron (MLP)

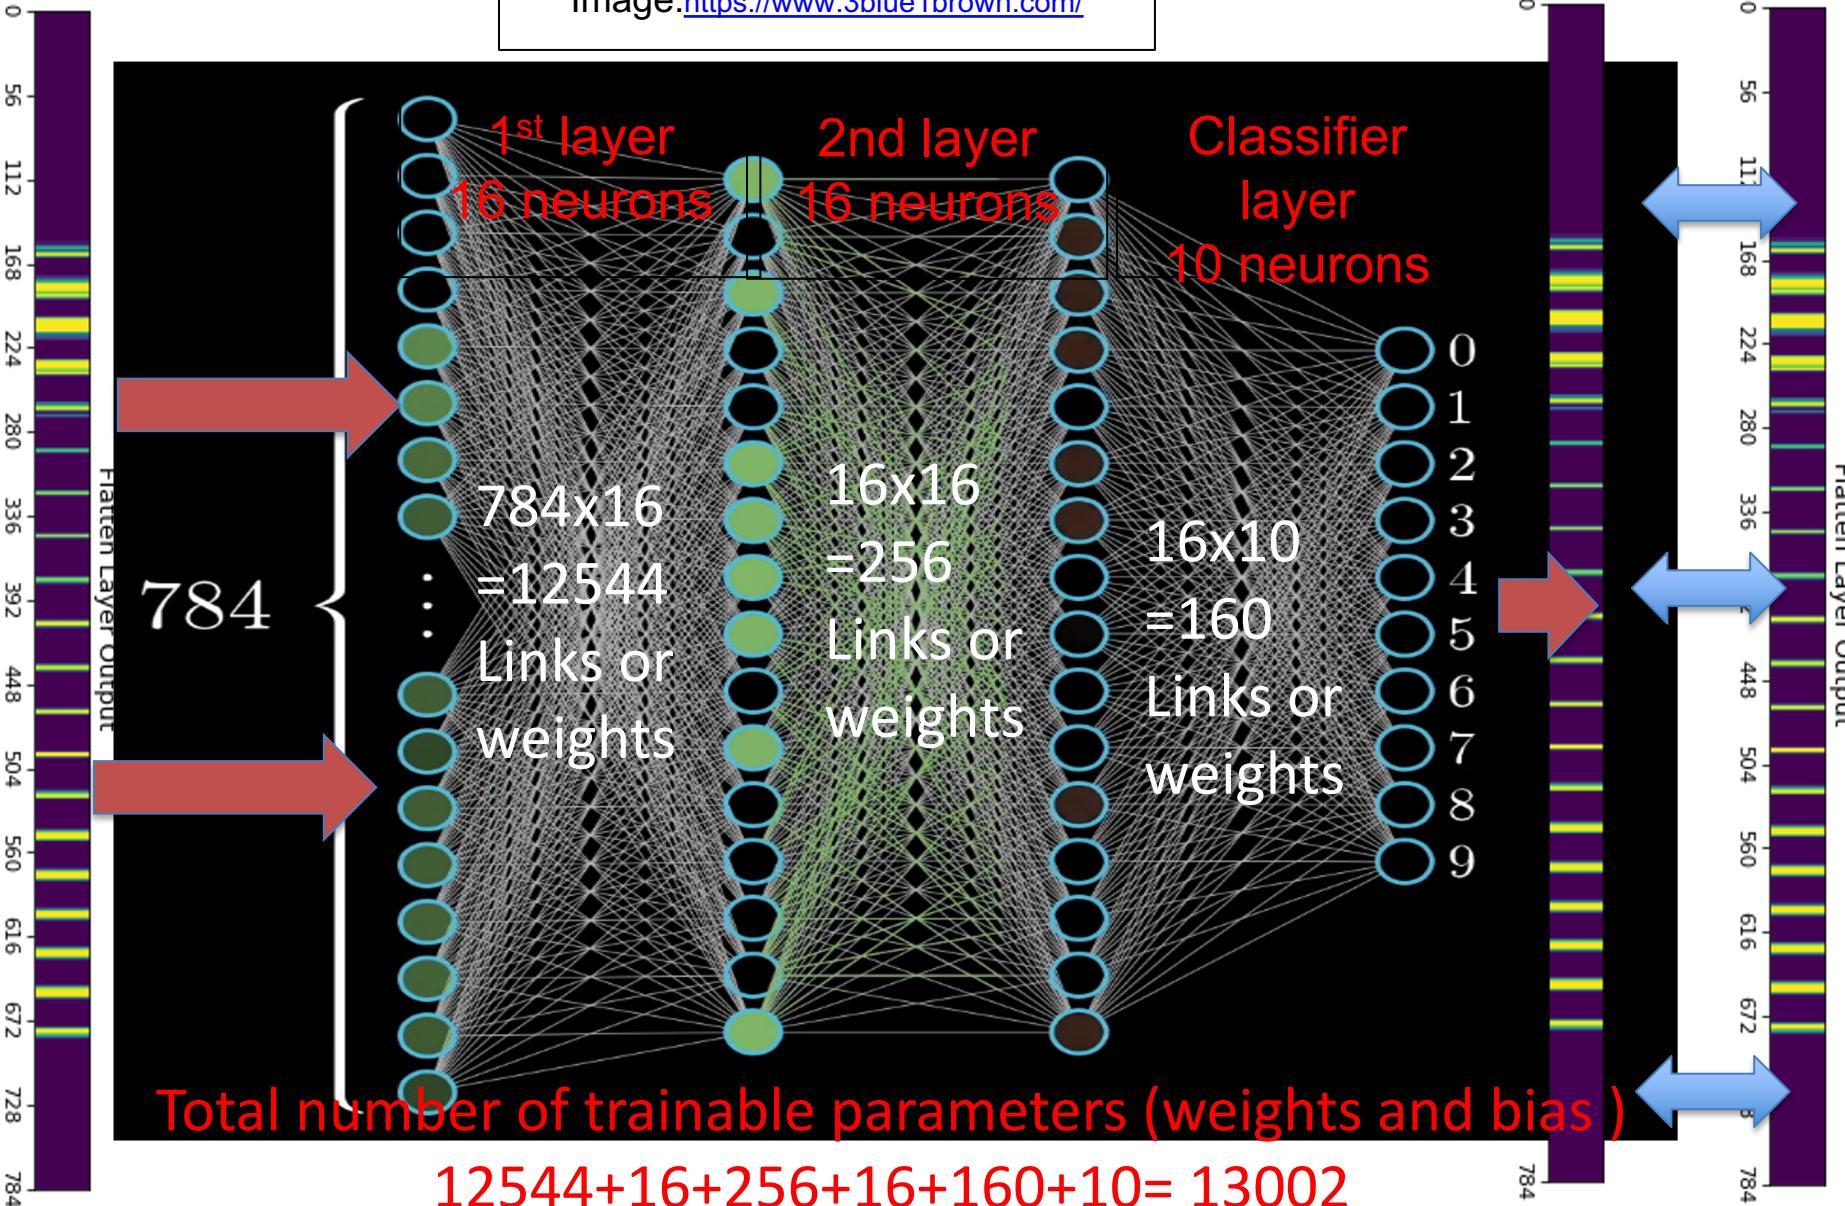
Simple MNIST MLP Neural Network

Input

output

labels

Image:<https://www.3blue1brown.com/>



MNIST Tensorflow code

```
import tensorflow as tf
print(tf.__version__)
import datetime, os

%load_ext tensorboard

mnist = tf.keras.datasets.mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(16, activation='relu'),
    tf.keras.layers.Dense(16, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='sgd', loss='SparseCategoricalCrossentropy',
metrics=['accuracy'])
model.summary()

logdir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
tensorboard_callback = tf.keras.callbacks.TensorBoard(logdir, histogram_freq=1)

model.fit(x_train, y_train, epochs=5, batch_size=10,
          validation_data=(x_test, y_test), callbacks=[tensorboard_callback])

model.evaluate(x_test, y_test, verbose=2)

%tensorboard --logdir logs
```

Show data download links Ignore outliers in chart scaling

Tooltip sorting method: default ▾

Smoothing



Horizontal Axis

 STEP RELATIVE WALL

Runs

Write a regex to filter runs

- 20210928-194736/train
- 20210928-194736/validation

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
dense (Dense)	(None, 16)	12560
dense_1 (Dense)	(None, 16)	272
dense_2 (Dense)	(None, 10)	170

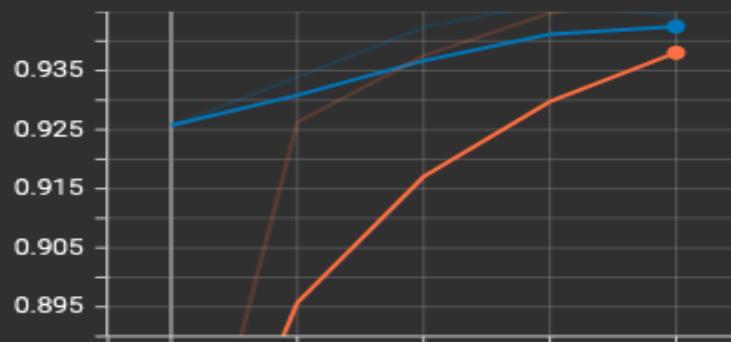
Total params: 13,002

Trainable params: 13,002

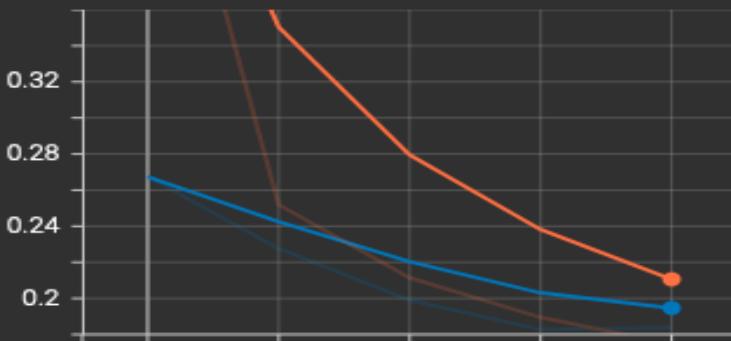
Non-trainable params: 0

Filter tags (regular expressions supported)

epoch_accuracy

epoch_accuracy
tag: epoch_accuracy

epoch_loss

epoch_loss
tag: epoch_loss

Exercise

- a) A neural network is used to classify a dataset composed of 5000 grayscale images of 0 to 9 digits. Each of the image has 32 x 32 pixels. The neural network contains 4 layers. The first layer is the input layer, the second layer has 20 neurons, the third one has 15 neurons, the last one is the softmax neuron of 10 classes. Please specify clearly the dimension of the matrices of each layer (both input weights and output) of the training process given a batch size of 50?

Layer one : input vector

Layer two:

input matrix and

weight matrix,

output matrix

Layer three:

input matrix

weight matrix,

output matrix

Layer four:

input matrix,

weight matrix, and

output

- b) What is the total number of trainable parameters?
- c) How many batch iterations are needed to finish one epoch ?

Exercise : Answer

- a) A neural network is used to classify a dataset composed of 5000 grayscale images of 0 to 9 digits. Each of the image has 32 x 32 pixels. The neural network contains 4 layers. The first layer is the input layer, the second layer has 20 neurons, the third one has 15 neurons, the last one is the softmax neuron of 10 classes. Please specify clearly the dimension of the matrices of each layer (both input weights and output) of the training process given a batch size of 50?

Layer one : input vector : $32 \times 32 = 1024$

Layer two:

input matrix = 50×1024

weight matrix = 1024×20

output matrix = 50×20

Layer three:

input matrix = 50×20

weight matrix = 20×15

output matrix = 50×15

Layer four:

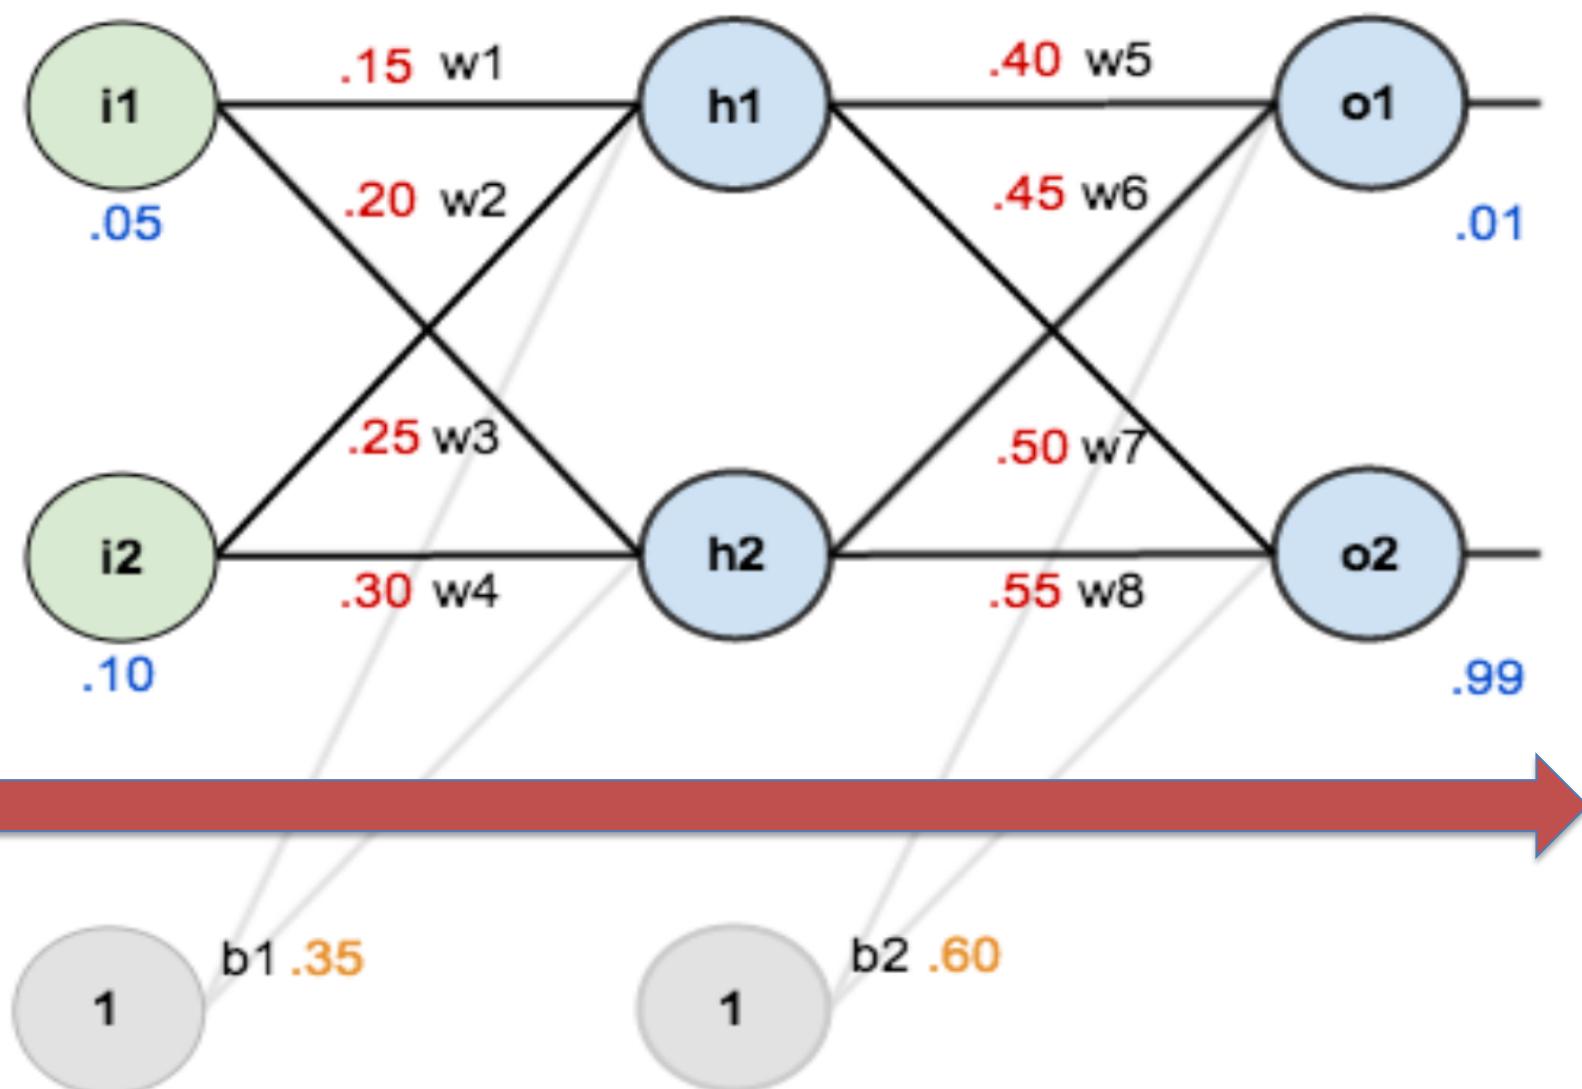
input matrix = 50×15

weight matrix = 15×10

output = 50×10

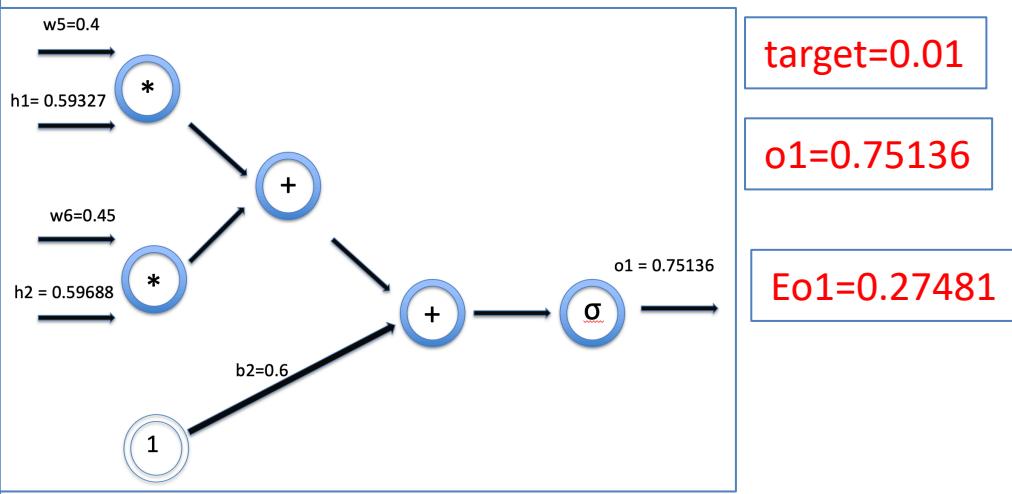
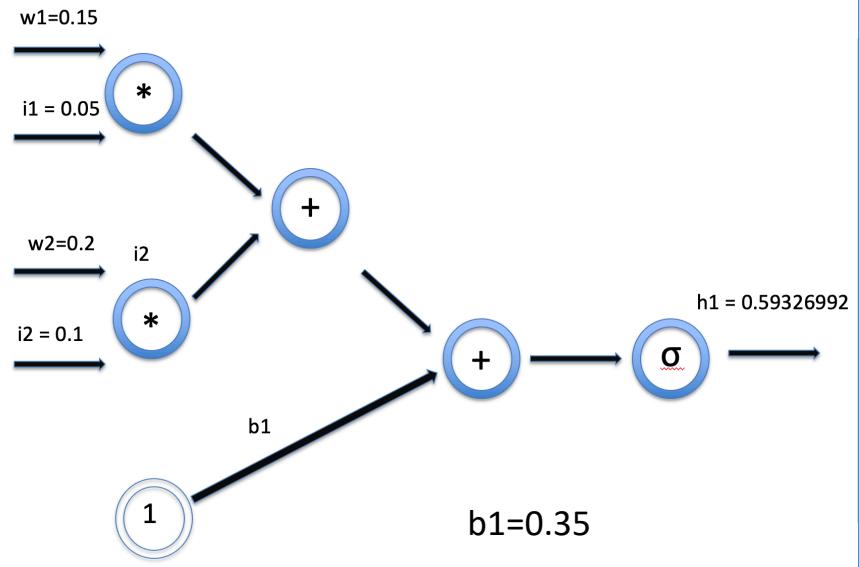
- b) Total number of trainable parameters = $(1024 \times 20 + 20) + (20 \times 15 + 15) + (15 \times 10 + 10) = 20975$
- c) How many batch iterations are needed to finish one epoch ? = $5000 / 50 = 100$

Multilayer Perceptron : Forward Path Calculation



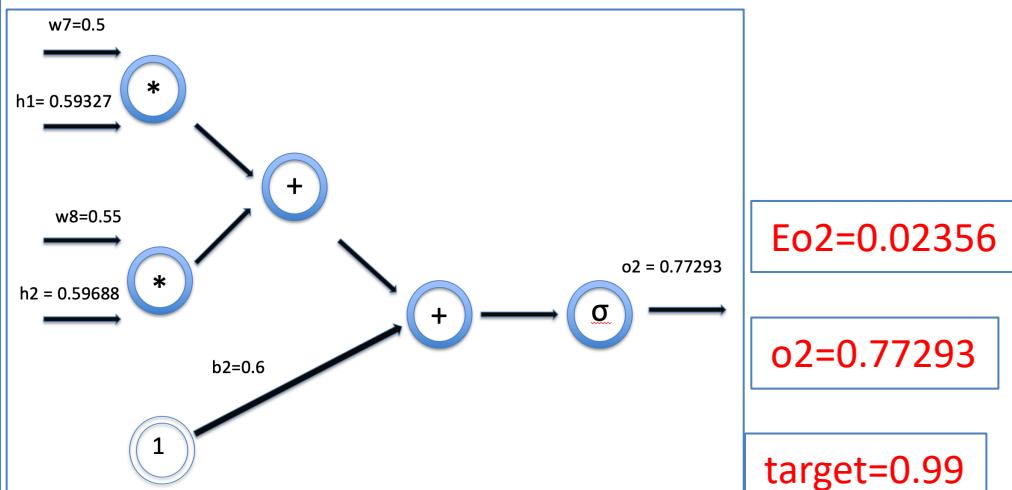
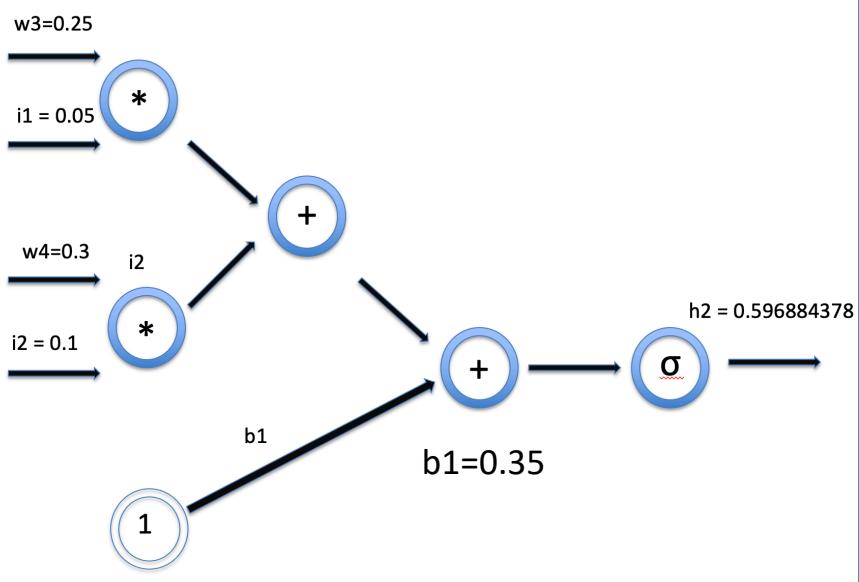
$$E_{o1} = \frac{1}{2}(target_{o1} - out_{o1})^2 = \frac{1}{2}(0.01 - 0.75136507)^2 = 0.274811083$$

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \tilde{y}_i)^2$$



$$E_{total} = \sum \frac{1}{2}(target - output)^2$$

$$E_{total} = E_{o1} + E_{o2} = 0.298371109$$



$$E_{o2} = 0.023560026$$

target=0.01

$o_1=0.75136$

$E_{o1}=0.27481$

$E_{o2}=0.02356$

$o_2=0.77293$

target=0.99

Python Code

Forward path calculation

```
# e constant
e = 2.7182818284
# initial values
i1 = 0.05
i2 = 0.10
# initial weights
w1 = 0.15
w2 = 0.20
w3 = 0.25
w4 = 0.30
w5 = 0.40
w6 = 0.45
w7 = 0.50
w8 = 0.55
# bias
b1 = 0.35
b2 = 0.60
# targets
To1 = 0.01
To2 = 0.99
```

```
# forward propagation
h1 = 1/(1+e**(-(w1*i1 + w2*i2+b1)))
print("h1: " + str(h1))

h2 = 1/(1+e**(-(w3*i1 + w4*i2+b1)))
print("h2: " + str(h2))

o1 = 1/(1+e**(-(w5*h1 + w6*h2+b2)))
print("o1: " + str(o1))

o2 = 1/(1+e**(-(w7*h1 + w8*h2+b2)))
print("o2: " + str(o2))
```

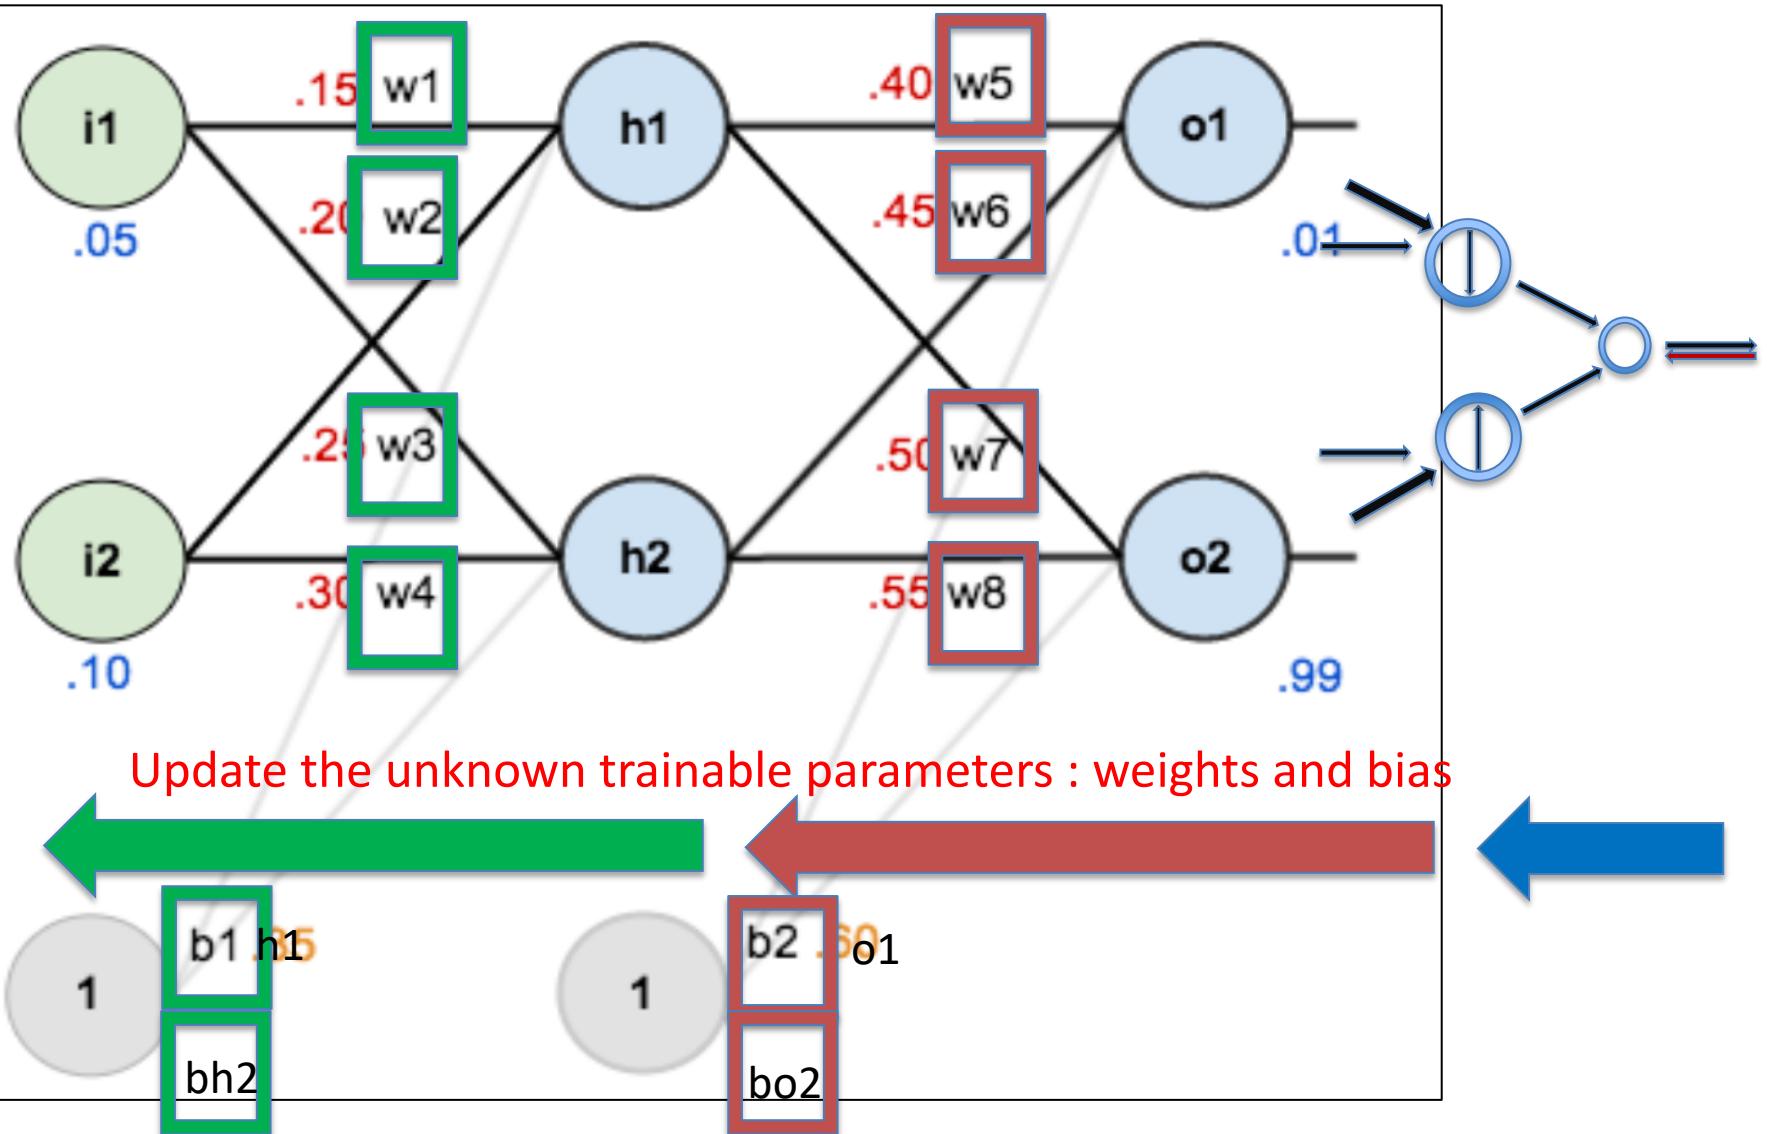
```
h1: 0.5932699921052087 h2: 0.5968843782577157
o1: 0.7513650695475076 o2: 0.772928465316421
```

```
# Error
Eo1 = 0.5*(To1-o1)**2
print("Error o1: " + str(Eo1))
Eo2 = 0.5*(To2-o2)**2
print("Error o2: " + str(Eo2))

E = Eo1 + Eo2
print("Total Error: " + str(E))
```

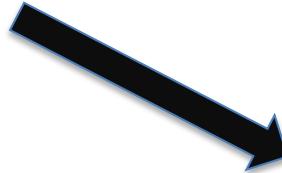
```
Error o1: 0.2748110831725904
Error o2: 0.023560025584942117
Total Error: 0.29837110875753253
```

Multilayer Perceptron : Backward Path Calculation



Backward Path from Objective Function

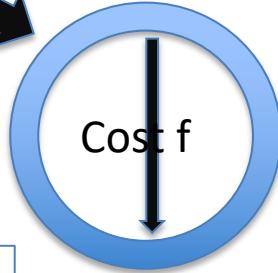
target1 = 0.01



$o_1 = 0.75136$



$\text{grad } o_1 = 0.74136$



composite diff : $\text{grad } o_1 = (\text{grad up}) * (\text{local grad})$

local grad = $(dE_1/do_1) = (o_1 - \text{target1})$

$\text{grad } o_1 = (1) (0.75136 - 0.01) = 0.74136$

$E_{o1} = \frac{1}{2} * (\text{target1} - o_1)^2$

$E_{o1} = 0.27481$

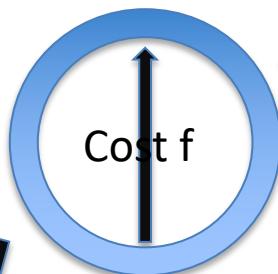
$ET = 0.298371$

$ET = E_1 + E_2$

$o_2 = 0.77293$



$\text{grad } o_2 = -0.21707$



composite diff : $\text{grad } o_2 = (\text{grad up}) * (\text{local grad})$

local grad = $(dE_2/do_2) = (o_2 - \text{target2})$

$E_{o2} = \frac{1}{2} * (\text{target2} - o_2)^2$

$ET = 0.298371$

$\text{grad } ET = 1.0$

target2 = 0.99

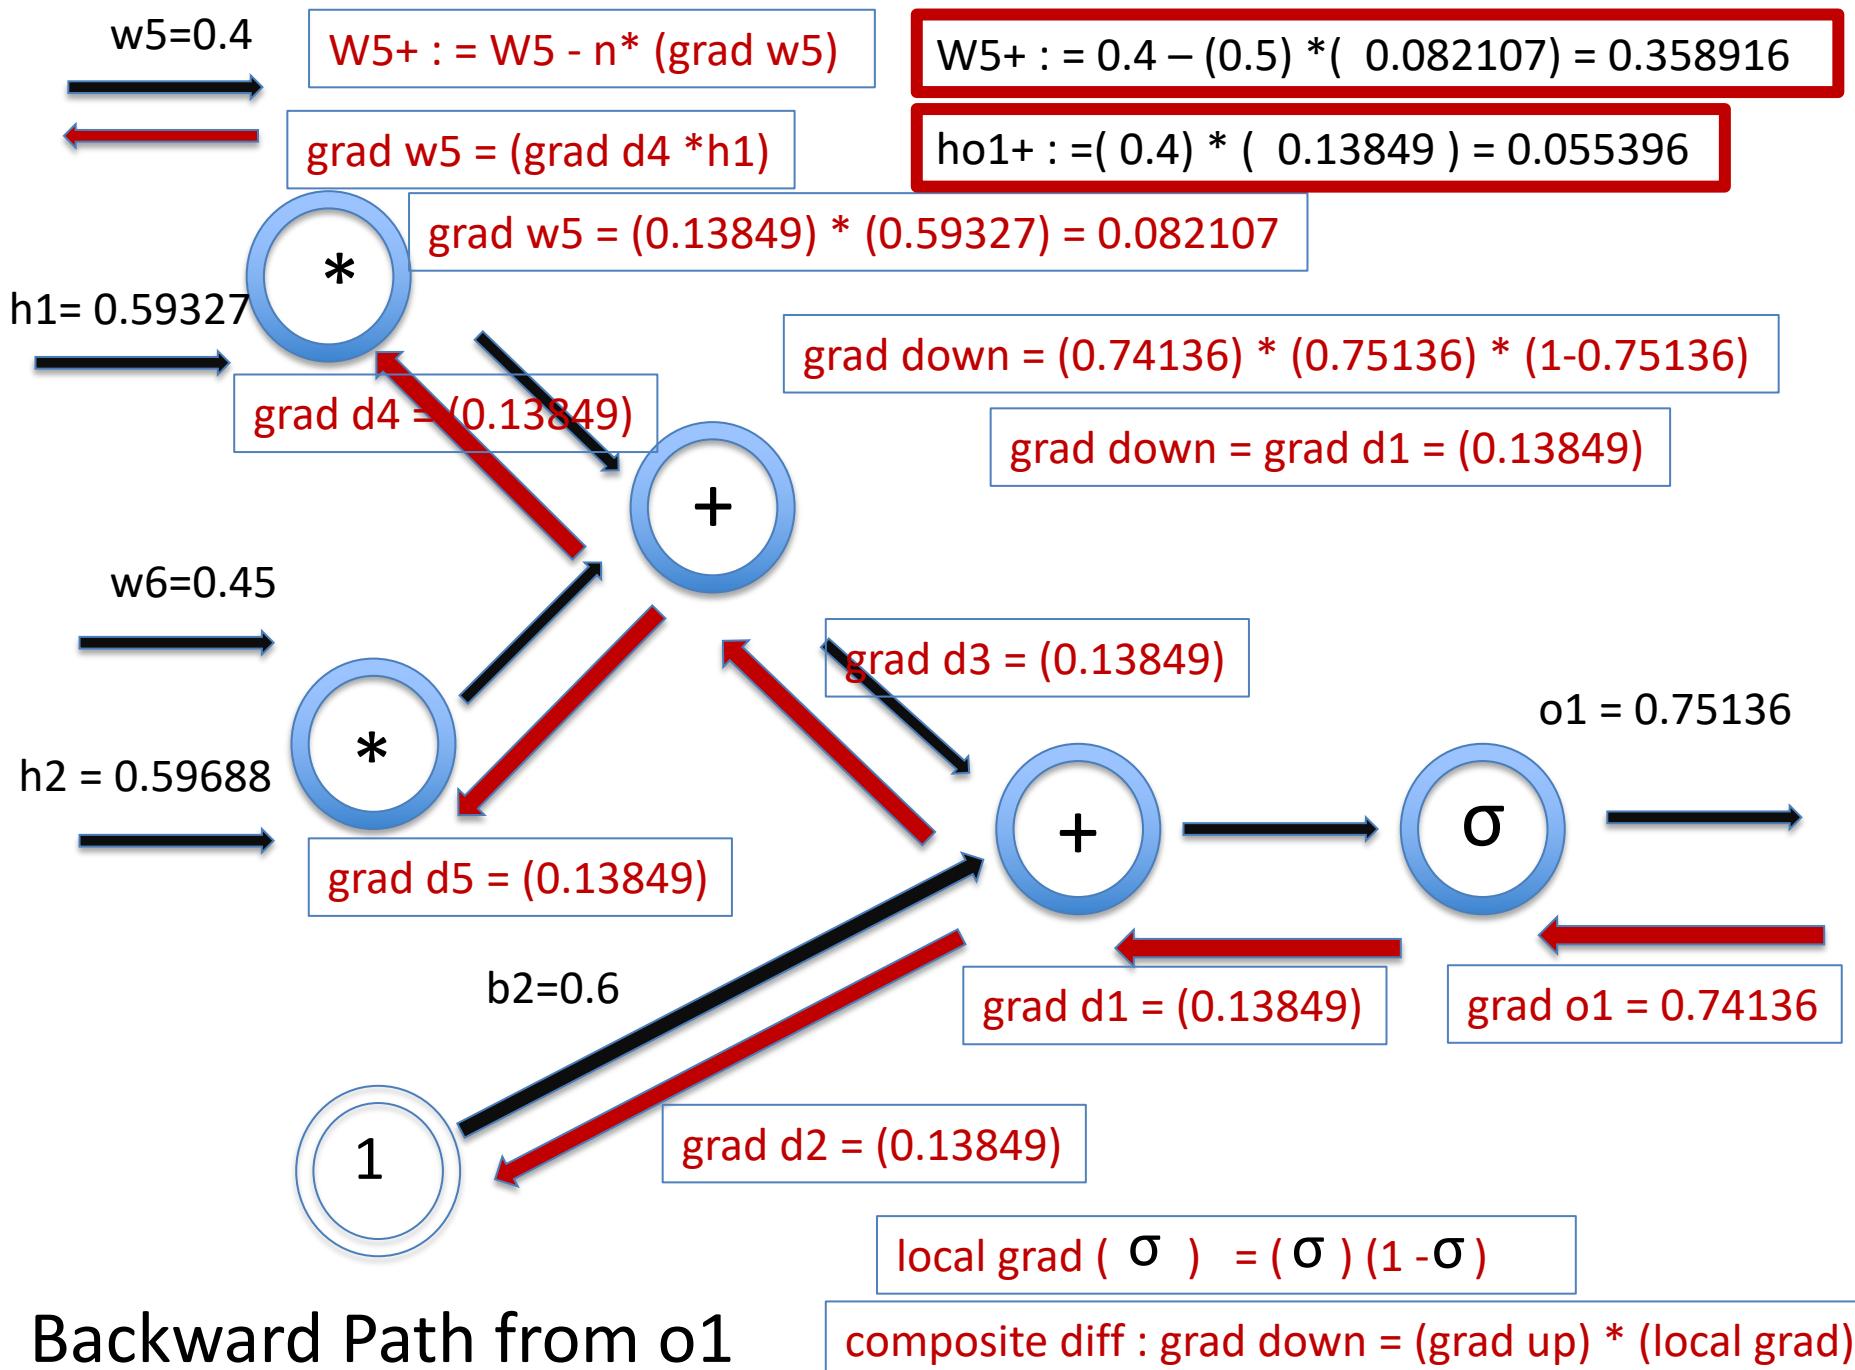
composite diff : $\text{grad } o_2 = (\text{grad up}) * (\text{local grad})$

$E_{o2} = 0.02356$

$\text{grad } E_{o2} = 1.0$

$\text{grad } E_{o1} = 1.0$

+



$w7=0.5$

$W7+ := W7 - n * (\text{grad } w7)$

$W7+ := 0.5 - (0.5) * (-0.022601) = 0.51130$

$\text{grad } w7 = (\text{grad } d9 * h1)$

$h02+ := (0.5) * (-0.038097) = -0.019485$

 $*$

$\text{grad } w7 = (-0.038097) * (0.59327) = -0.022601$

$n=0.5$

$h1 = 0.59327$

$\text{grad } d9 = (-0.038097)$

$\text{grad down} = (-0.21707) * (0.77293) * (1-0.77293)$

$\text{grad down} = \text{grad } d6 = (-0.038097)$

$w8=0.55$

 $+$

$\text{grad } d8 = (-0.038097)$

$o2 = 0.77293$

$h2 = 0.59688$

$\text{grad } d10 = (-0.038097)$

 $+$ σ $*$

$b2=0.6$

$\text{grad } d6 = (-0.038097)$

$\text{grad } o2 = -0.21707$

 1

$\text{grad } d7 = (-0.038097)$

$\text{local grad } (\sigma) = (\sigma)(1-\sigma)$

Backward Path from
 $o2$

$\text{composite diff : grad down} = (\text{grad up}) * (\text{local grad})$

$w1=0.15$

$W1+ := W1 - n * (\text{grad } w1)$

$\text{grad } w1 = (\text{grad } d14 * i1)$

$W1+ := 0.15 - (0.5) * (0.0004332) = 0.149783$

$i1 = 0.05$

$\text{grad } w1 = (0.008665) * (0.05) = 0.00043326$

$\text{grad } d14 = (0.008665)$

$\text{grad down} = (0.035911) * (0.59327) * (1 - 0.59327)$

$\text{grad down} = \text{grad } d1 = (0.008665)$

$w2=0.20$

$i2 = 0.1$

Backward Path from $h1$ to $w1$

$h1 = 0.59327$

$\text{grad } d15 = (0.008665)$

$\text{grad } d13 = (0.008665)$

$+$

σ

$b1=0.35$

$\text{grad } d11 = (0.008665)$

$\text{grad } h1 = 0.035911$

1

$\text{grad } d12 = (0.008665)$

$\text{grad } h1-o1 := (0.4) * (0.13849) = 0.055396$

$\text{grad } h1-o2 := (0.5) * (-0.0038097) = -0.019485$

$\text{local grad } (\sigma) = (\sigma)(1-\sigma)$

$\text{grad } h1 := (\text{grad } h1-o1) + (\text{grad } h1-o2) = (0.055396 - 0.019485)$

$\text{composite diff : grad down} = (\text{grad up}) * (\text{local grad})$

```

# Backpropagation
n = 0.5

grad_o1 = o1 - To1
grad_o2 = o2 - To2

print("grad_o1: "+str(grad_o1))
print("grad_o2: "+str(grad_o2))

grad_d1 = (grad_o1)*o1*(1-o1)
grad_d2 = (grad_o2)*o2*(1-o2)

print("grad_d1: "+str(grad_d1))
print("grad_d2: "+str(grad_d2))

grad_w5 = grad_d1*h1
print("grad_w5: "+str(grad_w5))

w5_final = w5 - n*grad_w5

grad_w6 = grad_d1*h2
print("grad_w6: "+str(grad_w6))
w6_final = w6 - n*grad_w6

grad_w7 = grad_d2*h1
print("grad_w7: "+str(grad_w7))
w7_final = w7 - n*grad_w7

grad_w8 = grad_d2*h2
print("grad_w8: "+str(grad_w8))
w8_final = w8 - n*grad_w8

```

```

grad_h1 = w5*grad_d1 + w7*grad_d2
print("grad_h1: "+str(grad_h1))
grad_d11 = grad_h1*h1*(1-h1)
print("grad_d11: "+str(grad_d11))

grad_w1 = grad_d11*i1
print("grad_w1: "+str(grad_w1))
w1_final = w1 - n*grad_w1

grad_w2 = grad_d11*i1
print("grad_w2: "+str(grad_w2))
w2_final = w2 - n*grad_w2

grad_h2 = w6*grad_d1 + w8*grad_d2
print("grad_h2: "+str(grad_h2))
grad_d22 = grad_h2*h2*(1-h2)
print("grad_d22: "+str(grad_d22))

grad_w3 = grad_d22*i1
print("grad_w3: "+str(grad_w3))
w3_final = w3 - n*grad_w3

grad_w4 = grad_d22*i2
print("grad_w4: "+str(grad_w4))
w4_final = w4 - n*grad_w4

```

```
print("w1+: "+str(w1_final))
print("w2+: "+str(w2_final))
print("w3+: "+str(w3_final))
print("w4+: "+str(w4_final))
print("w5+: "+str(w5_final))
print("w6+: "+str(w6_final))
print("w7+: "+str(w7_final))
print("w8+: "+str(w8_final))
```

w1+: 0.1497807161327648
w2+: 0.19978071613276482
w3+: 0.24975114363237164
w4+: 0.29950228726474326
w5+: 0.35891647971775653
w6+: 0.4086661860761087
w7+: 0.5113012702391395
w8+: 0.5613701211083925

grad_o1: 0.7413650695475076
grad_o2: -0.21707153468357898
grad_d1: 0.13849856162945076
grad_d2: -0.03809823651803844
grad_w5: 0.08216704056448701
grad_w6: 0.08266762784778263
grad_w7: -0.02260254047827904
grad_w8: -0.02274024221678477
grad_h1: 0.036350306392761086
grad_d11: 0.00877135468940779
grad_w1: 0.0004385677344703895
grad_w2: 0.0004385677344703895
grad_h2: 0.04137032264833171
grad_d22: 0.009954254705134271
grad_w3: 0.0004977127352567136
grad_w4: 0.0009954254705134271

```

# Tensorflow code for the example
import tensorflow as tf
import keras

# sets the input [i1, i2]
x = tf.constant([0.05, 0.1], shape=[1,2],
dtype=tf.float32)

#setting the weight matrix for first fully
connected layer [w1, w3];[w2, w4]
w1 = tf.constant([[0.15,0.25],[0.2,0.3]])

#setting the weight matrix for second fc layer
[w5, w7];[w6, w8]
w2 = tf.constant([[0.4,0.5],[0.45,0.55]])

# sets the target matrix [t1, t2]
t = tf.constant([0.01,0.99])

#sets the bias1.
b1 = tf.constant([0.35, 0.35])

#sets the bias2.
b2 = tf.constant([0.6, 0.6])

def forward(x,w1,w2,b1,b2):
    sum1=tf.matmul(x,w1)+b1
    H = tf.nn.sigmoid(sum1)
    sum2=tf.matmul(H,w2)+b2
    O = tf.nn.sigmoid(sum2)
    return O

```

```

def backward(x,w1,w2,b1,b2,t):
    with tf.GradientTape() as g:
        g.watch([w1,w2,x,b1,b2])
        O = forward(x, w1, w2, b1, b2)
        Etotal = tf.losses.mse(t, O)
        print("Total loss=" + str(Etotal))

#calculates the gradients for each part
dEtotal_dw = g.gradient(Etotal,
[w1,w2,x,b1,b2])
dEtotal_dw1, dEtotal_dw2,
dEtotal_dx, dEtotal_db1, dEtotal_db2
= dEtotal_dw
return dEtotal_dw1, dEtotal_dw2,
dEtotal_dx,dEtotal_db1,dEtotal_db2

dEtotal_dw1,dEtotal_dw2,dEtotal_dx,dEtotal_db1,
dEtotal_db2=backward(x,w1,w2,b1,b2,t)

#the learning rate is set to be 0.5
w1 = w1 - 0.5*dEtotal_dw1
w2 = w2 - 0.5*dEtotal_dw2

b1 = b1 - 0.5*dEtotal_db1
b2 = b2 - 0.5*dEtotal_db2

print('The new w1 is',w1.numpy())
print('The new w2 is',w2.numpy())
print('The new b1 is',b1.numpy())
print('The new b2 is',b2.numpy())

```

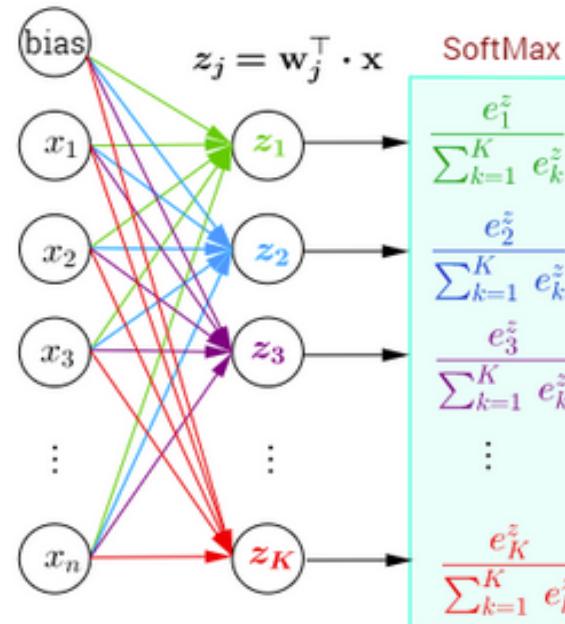
Total loss = tf.Tensor([0.2983711], shape=(1,), dtype=float32)
The new w1 is [[0.14978072 0.24975115] [0.19956143 0.2995023]]
The new w2 is [[0.3589165 0.5113013] [0.40866616 0.56137013]]
The new b1 is [0.3456143 0.34502286] The new b2 is [0.53075075 0.61904913]

Multi-Class Classification with NN and SoftMax Function

$$S(z_i) = \frac{e^{z_i}}{\sum_{j=1}^N e^{z_j}}$$

$$\mathbf{z} = \begin{bmatrix} z_1 \\ z_2 \\ z_3 \\ \vdots \\ z_K \end{bmatrix} = \begin{bmatrix} \mathbf{w}_1^\top \\ \mathbf{w}_2^\top \\ \mathbf{w}_3^\top \\ \vdots \\ \mathbf{w}_K^\top \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{bmatrix}$$

Softmax Function



Categorical (Softmax) Cross Entropy Loss (Statistical Learning)

s

Softmax

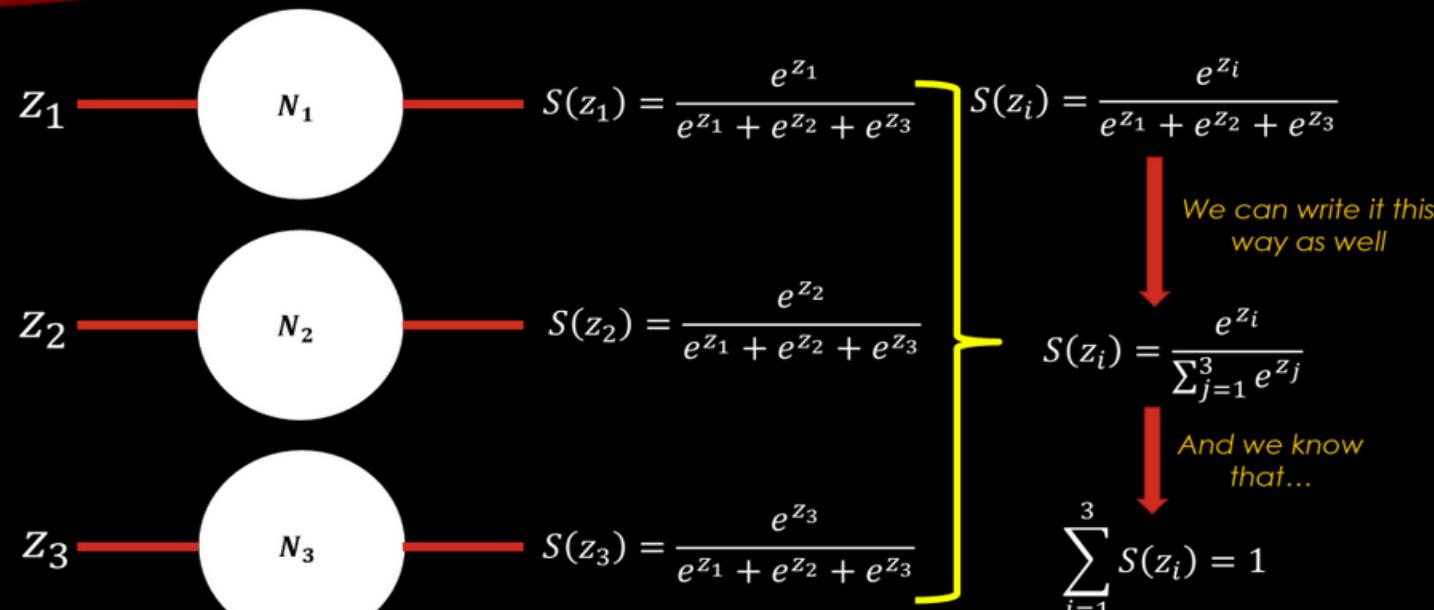
Cross-Entropy
Loss

$$f(s)_i = \frac{e^{s_i}}{\sum_j^C e^{s_j}}$$

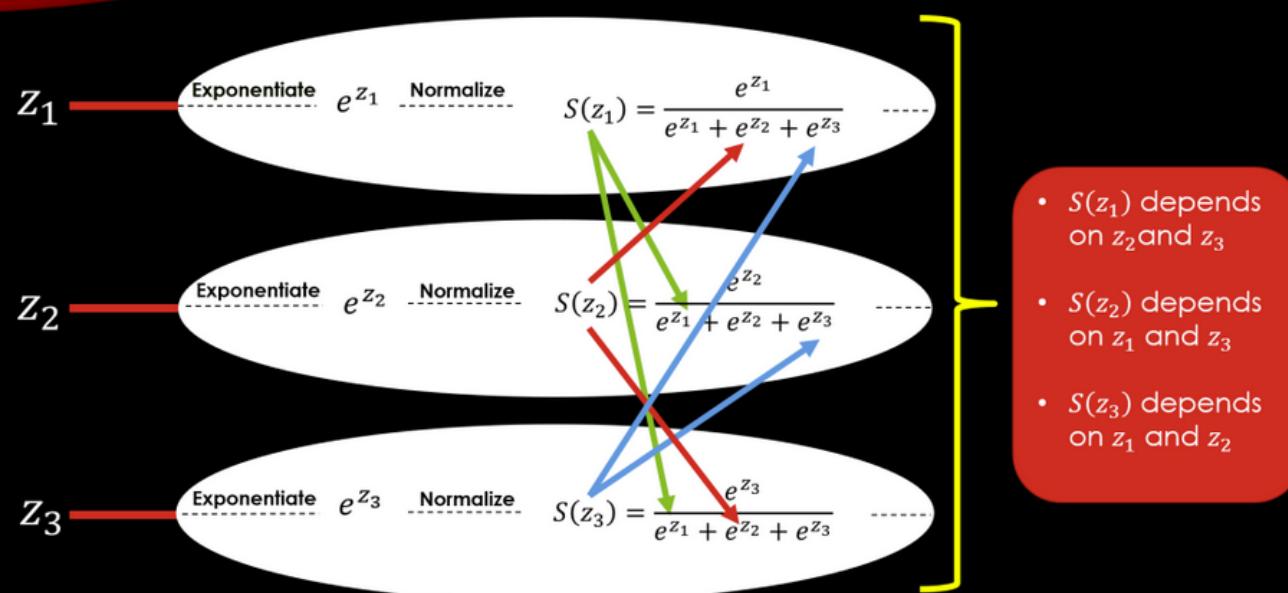
$$CE = - \sum_i^C t_i \log(f(s)_i)$$

$$\frac{\partial E}{\partial z} = \frac{\partial (\text{cross entropy})}{\partial (z = \mathbf{w}^\top \mathbf{x} + b)} = \frac{\partial (y_i \ln f(s)_i)}{\partial z} = (o_i - y_i)$$

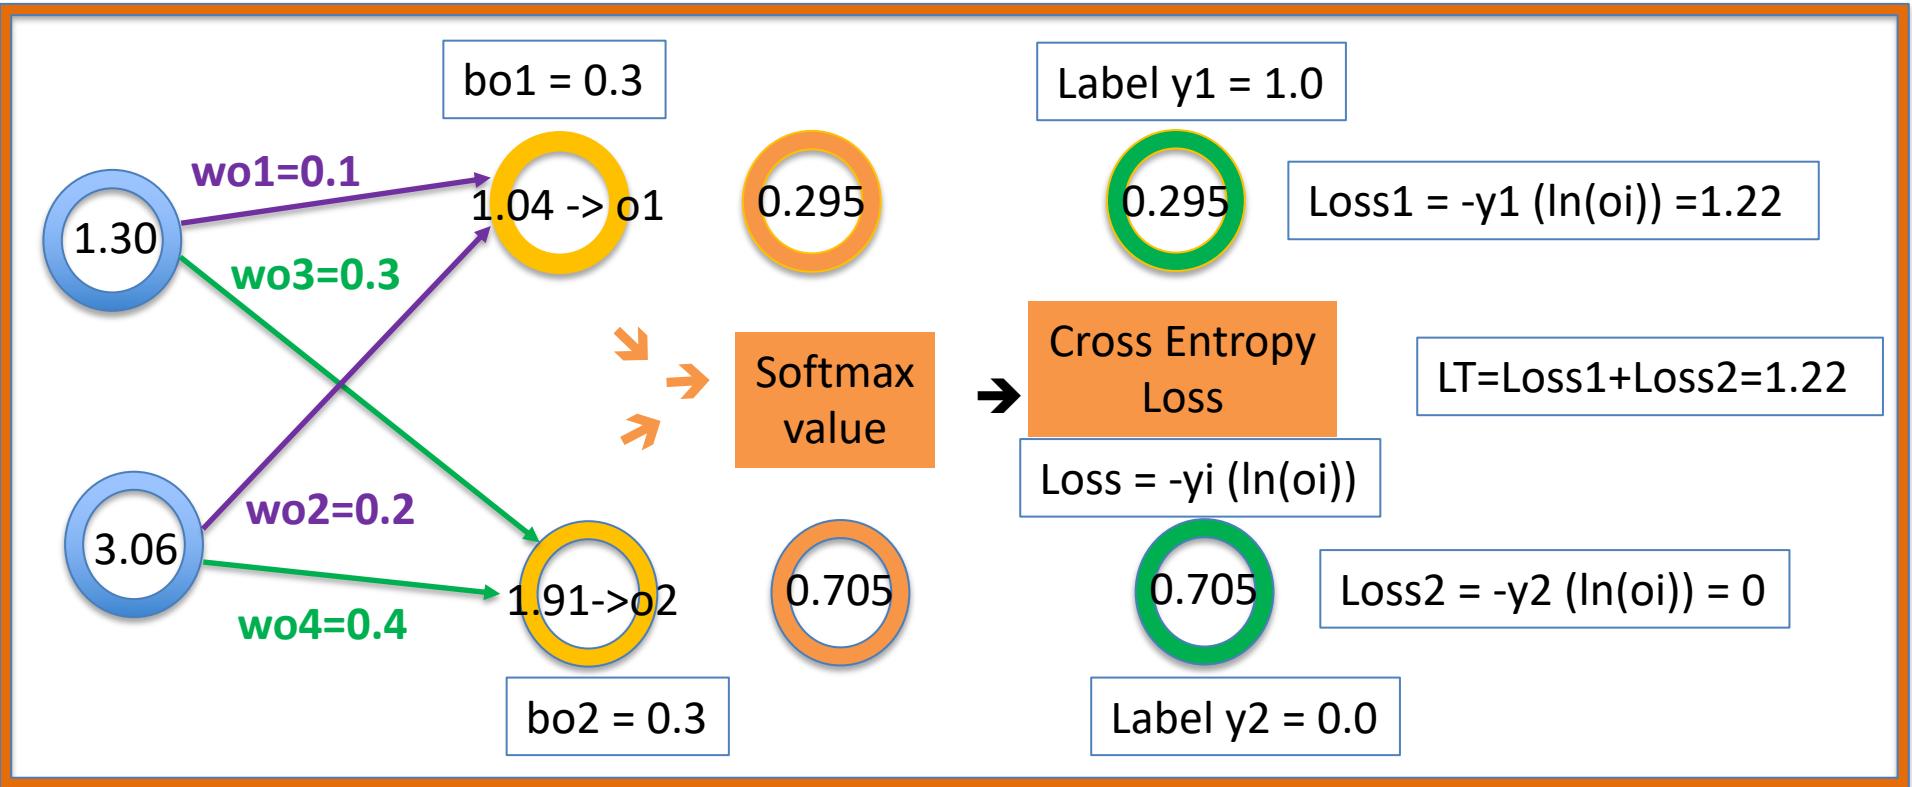
A SOFTMAX UNIT: $S(z_i)$



UNDERSTANDING THE DEPENDENCIES IN SOFTMAX



<https://www.mldawn.com/the-derivative-of-softmaxz-function-w-r-t-z/>



CNN : Backward Path (derivatives)

Learning rate
= $n = 0.5$

$$\frac{\partial S(z_i)}{\partial z_j} = \begin{cases} S(z_i) \times (1 - S(z_i)) & \text{if } i = j \\ -S(z_i) \times S(z_j) & \text{if } i \neq j \end{cases}$$

$$\frac{\partial E}{\partial Z_k} = O_k - y_k$$

weights
bias

$$w+ = w - n * (dE/dw)$$

$$dE/dw = (dE/dS) * (dS/dw)$$

Softmax
value

dL/dS

Cross Entropy
Loss (L)

$dL/dL = 1$

Exercises : Represent the following NN model as a computational graph and build a python code to update w and b

Input

$x_1 = -0.04$

$x_2 = -0.42$

NN Model

$b_1 = -1.6$

$b_2 = 0.7$

Output

$S(z)$
 $b_3 = 0$

$S(z)$
 $b_4 = 0$

$S(z)$
 $b_5 = 1$

Loss

CCE

CCE

CCE

Target

$T = 0$

$T = 1$

$T = 0$

Categorical (Softmax) Cross Entropy Loss

$$w_1 = -2.5$$

$$w_2 = -1.5$$

$$w_3 = 0.6$$

$$w_4 = 0.4$$

$$w_5 = -0.1$$

$$w_6 = 2.4$$

$$w_7 = -2.2$$

$$w_8 = -1.5$$

$$w_9 = -5.2$$

$$w_{10} = 3.7$$

```
# the code for updating the parameters in question 2
import tensorflow as tf
import keras

x = tf.constant([-0.04, -0.42], shape=[1, 2], dtype=tf.float32) # sets the input [X1, X2]

w1 = tf.constant([[-2.5, -1.5], [0.6, 0.4]]) #setting the weight matrix for first fully connected layer [w1, w2];[w3, w4]
w2 = tf.constant([[-0.1, 2.4, -2.2], [-1.5, -5.2, 3.7]]) #setting the weight matrix for second fc layer [w5, w6, w7];[w8, w9, w10]

t = tf.constant([[0, 1, 0]], dtype=tf.float32) # sets the target matrix [t1 t2,t3]

b1 = tf.constant([-1.6, 0.7]) #sets the bias1.
b2 = tf.constant([0, 0, 1], dtype=tf.float32) #sets the bias2.

def forward(x,w1,w2,b1,b2):
    sum1=tf.matmul(x,w1)+b1
    H = tf.nn.sigmoid(sum1)
    print(H)
    sum2=tf.matmul(H,w2)+b2
    O = tf.nn.softmax(sum2)
    return O

O = forward(x, w1, w2, b1, b2)

print('Forward path output is:', O.numpy())
```

```

def backward(x,w1,w2,b1,b2,t):
    with tf.GradientTape() as g:
        g.watch([w1,w2,x,b1,b2])
        O = forward(x, w1, w2, b1, b2)
        Etotal = tf.losses.categorical_crossentropy(t, O)
        print("Total loss = " , Etotal.numpy())
    dEtotal_dw = g.gradient(Etotal, [w1,w2,x,b1,b2]) #calculates the gradients
    dEtotal_dw1, dEtotal_dw2, dEtotal_dx, dEtotal_db1, dEtotal_db2 = dEtotal_dw
    return dEtotal_dw1,dEtotal_dw2,dEtotal_dx,dEtotal_db1,dEtotal_db2

dEtotal_dw1,dEtotal_dw2,dEtotal_dx,dEtotal_db1,dEtotal_db2 =
backward(x,w1,w2,b1,b2,t)

w1 = w1 - 0.5*dEtotal_dw1 #the learning rate is set to be 0.5
w2 = w2 - 0.5*dEtotal_dw2

b1 = b1 - 0.5*dEtotal_db1
b2 = b2 - 0.5*dEtotal_db2

print('The new w1 is:')
print(w1.numpy())
print('The new w2 is:')
print(w2.numpy())
print('The new b1 is:')
print(b1.numpy())
print('The new b2 is:')
print(b2.numpy())

```

`tf.Tensor([[0.14779511 0.64382386]], shape=(1, 2), dtype=float32)`
 Forward path output is: [[0.01729538 0.00231123. 0.9803934]]
`tf.Tensor([[0.14779511 0.64382386]], shape=(1, 2), dtype=float32)`
 Total loss = [6.069976]
 The new w1 is: [[-2.5114694 -1.4596888]
 [0.47957253 0.8232676]]
 The new w2 is: [[-0.10127809 2.4737267 -2.2724488]
 [-1.5055676 -4.878832 3.3843997]]
 The new b1 is: [-1.313268 -0.30777997]
 The new b2 is: [-0.00864769 0.4988444 0.5098033]

Exercise : write a simple python code for the example

Homework

Question

CIFAR stands for Canadian Institute For Advanced Research and 10 refers to 10 classes. The CIFIR 10 dataset consists of 60000 32x32 color images in 10 classes, with 6000 images per class. There are 50000 training images and 10000 test images.

Convert CIFIR-10 images to greyscale images or reshape the CIFIR-10 images to vectors similar to that of MNIST dataset.

Modify the simple MLP TensorFlow tutorial which uses the MNIST data to train the modified CIFIR 10 dataset. Plot the loss and accuracy of the training and testing results.

Helpful web links:

<https://www.tensorflow.org/tutorials/quickstart/beginner>

https://www.tensorflow.org/api_docs/python/tf/image/rgb_to_grayscale

<https://datascience.stackexchange.com/questions/24459/how-to-give-cifar-10-as-an-input-to-mlp>

```
import tensorflow as tf
print(tf.__version__)
import datetime, os

%load_ext tensorboard

cifar10 = tf.keras.datasets.cifar10 # mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = cifar10.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(32, 32, 3)),
    tf.keras.layers.Dense(16, activation='relu'),
    tf.keras.layers.Dense(16, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='sgd', loss='SparseCategoricalCrossentropy',
metrics=['accuracy'])
model.summary()

logdir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
tensorboard_callback = tf.keras.callbacks.TensorBoard(logdir, histogram_freq=1)

model.fit(x_train, y_train, epochs=5, batch_size=50, validation_data=(x_test,
y_test), callbacks=[tensorboard_callback])

model.evaluate(x_test, y_test, verbose=2)

%tensorboard --logdir logs
```

2.6.0

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>

170500096/170498071 [=====] - 2s 0us/step

170508288/170498071 [=====] - 2s 0us/step

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 3072)	0
dense (Dense)	(None, 16)	49168
dense_1 (Dense)	(None, 16)	272
dense_2 (Dense)	(None, 10)	170

Total params: 49,610

Trainable params: 49,610

Non-trainable params: 0

Epoch 1/5

1000/1000 [=====] - 4s 3ms/step - loss: 2.0560 - accuracy: 0.2415 - val_loss: 1.8881 - val_accuracy: 0.3139

Epoch 2/5

1000/1000 [=====] - 3s 3ms/step - loss: 1.8462 - accuracy: 0.3332 - val_loss: 1.7885 - val_accuracy: 0.3547

Epoch 3/5

1000/1000 [=====] - 3s 3ms/step - loss: 1.7825 - accuracy: 0.3550 - val_loss: 1.7762 - val_accuracy: 0.3538

Epoch 4/5

1000/1000 [=====] - 3s 3ms/step - loss: 1.7417 - accuracy: 0.3726 - val_loss: 1.7417 - val_accuracy: 0.3611

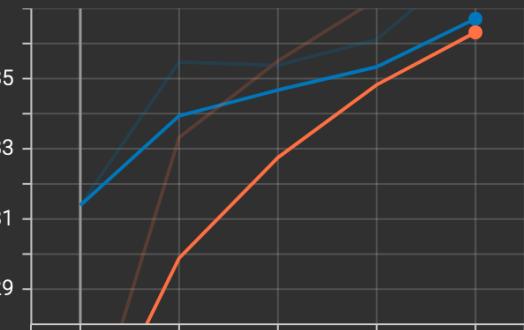
Epoch 5/5

1000/1000 [=====] - 3s 3ms/step - loss: 1.7157 - accuracy: 0.3827 - val_loss: 1.7162 - val_accuracy: 0.3849

313/313 - 1s - loss: 1.7162 - accuracy: 0.3849

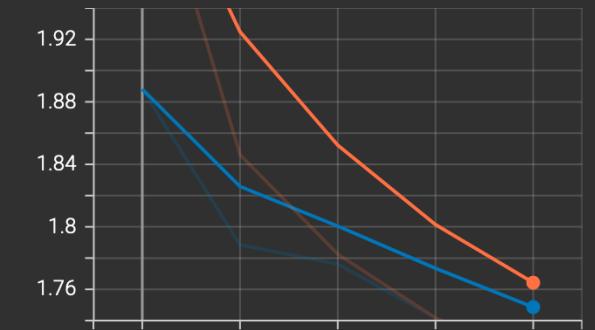
epoch_accuracy

epoch_accuracy
tag: epoch_accuracy



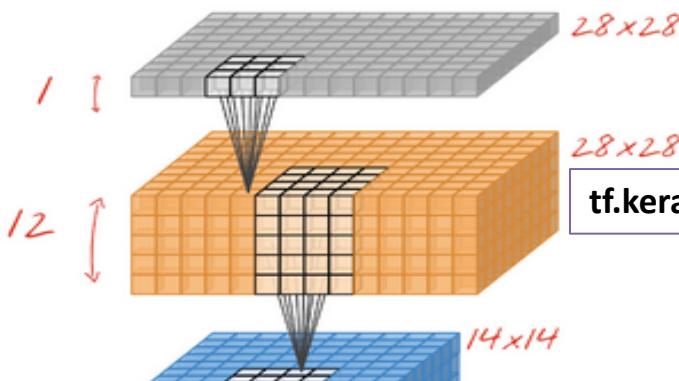
epoch_loss

epoch_loss
tag: epoch_loss



Convolution NN : MNIST 28x28 Pixels

[FW, FH, C, H] = [filter size ? X ?, input channel (Depth), output channel (no. of filter)]



`tf.keras.layers.Reshape(input_shape=(28*28,), target_shape=(28, 28, 1))`

*Convolutional 3x3 filters=12
W₁[3, 3, 1, 12]*

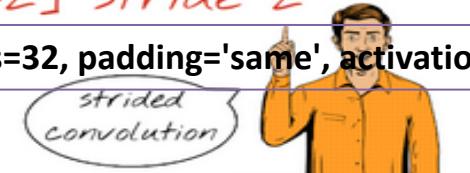
`tf.keras.layers.Conv2D(kernel_size=3, filters=12, padding='same', activation='relu')`

*Convolutional 6x6 filters=24
W₂[6, 6, 12, 24] stride 2*

`tf.keras.layers.Conv2D(kernel_size=6, filters=24, padding='same', activation='relu', strides=2)`

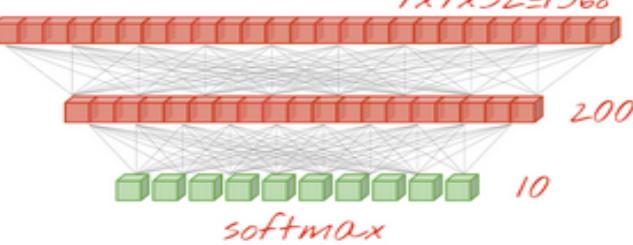
*Convolutional 6x6 filters=32
W₃[6, 6, 24, 32] stride 2*

`tf.keras.layers.Conv2D(kernel_size=6, filters=32, padding='same', activation='relu', strides=2)`



`tf.keras.layers.Flatten()`

<https://codelabs.developers.google.com/codelabs/cloud-tensorflow-mnist/#0>



Dense layer

W₄[1568, 200]

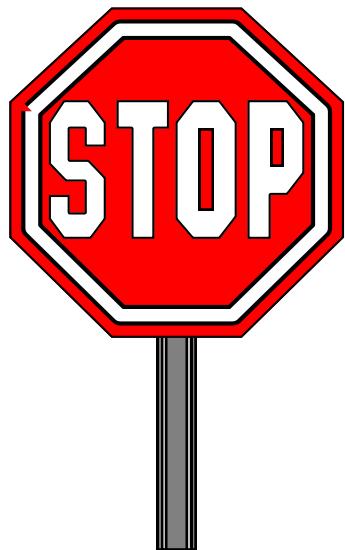
`tf.keras.layers.Dense(200, activation='relu')`

Softmax dense layer

W₅[200, 10]

`tf.keras.layers.Dense(10, activation='softmax')`

The End



- The End!

